

# Class 4

## Hadoop-Related Tools



CSCI 6830

BIG DATA ANALYTICS  
WITH HADOOP AND R

# Apache Tez

APPLICATION FRAMEWORK FOR A  
COMPLEX DIRECTED-ACYCLIC-GRAPH  
OF TASKS FOR PROCESSING DATA.






# Apache Tez




- Expressive dataflow definition APIs
- Flexible Input-Processor-Output runtime model
- Data types do not matter
- Simplifying deployment
- Performance gains over Map Reduce
- Optimal resource management
- Plan reconfiguration at runtime
- Dynamic physical data flow decisions

# Apache Tez

- Tez can be used to process data, that earlier took multiple MR jobs, now it is a single Tez job



Pig/Hive - MR



Pig/Hive - Tez





# Apache Tez



- Allows an application to be modeled as a data flow
- Edges (nodes) reflect the data movement
- Vertices represent the data processing tasks
- After processing the query, Pig or Hive generate a directed acyclic graph (DAG)
- DAG is ready for execution
- Tez is well suited as the executor for Pig or Hive jobs
- Tez is not only an engine though
- Tez also provides common modules to develop applications
- Tez focuses on flexibility and customizability.
- Developers may implement actual MapReduce jobs via the Tez library.

# Apache Tez Example

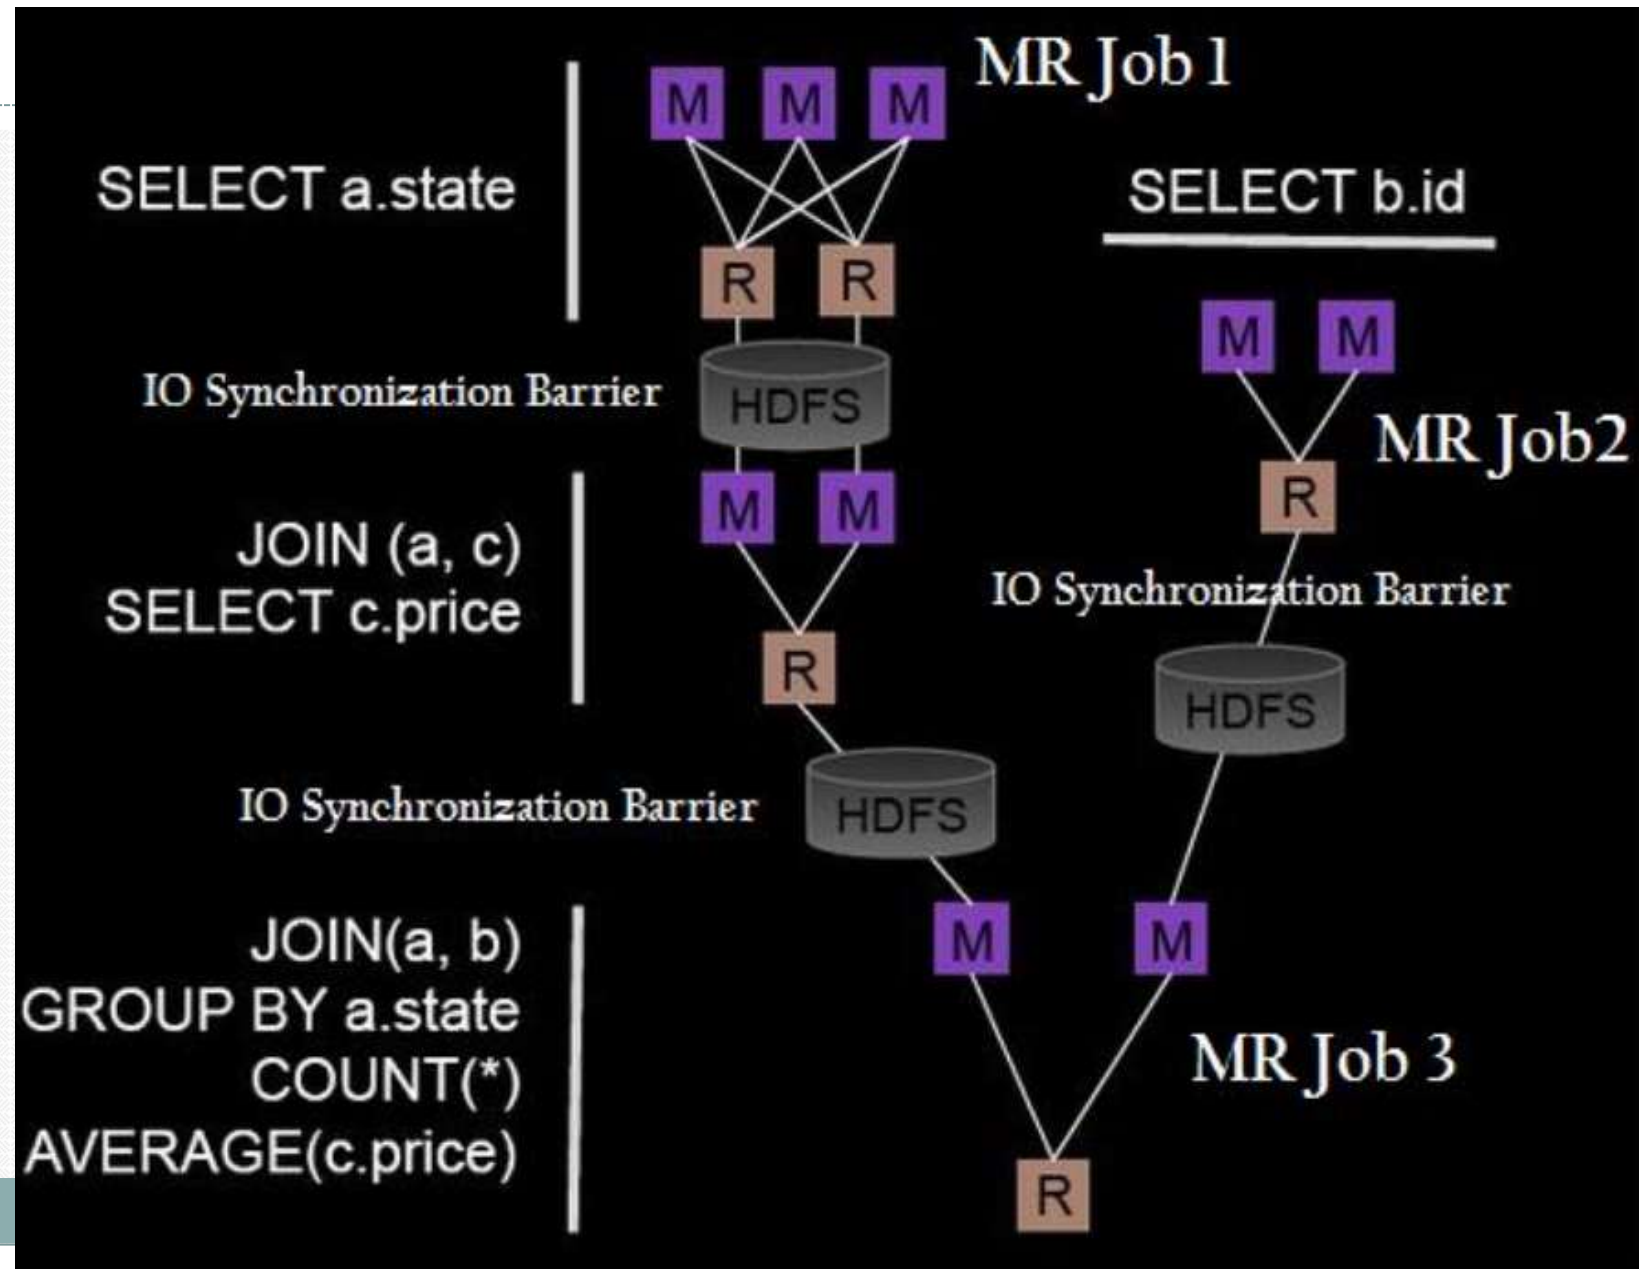


- Example Hive query

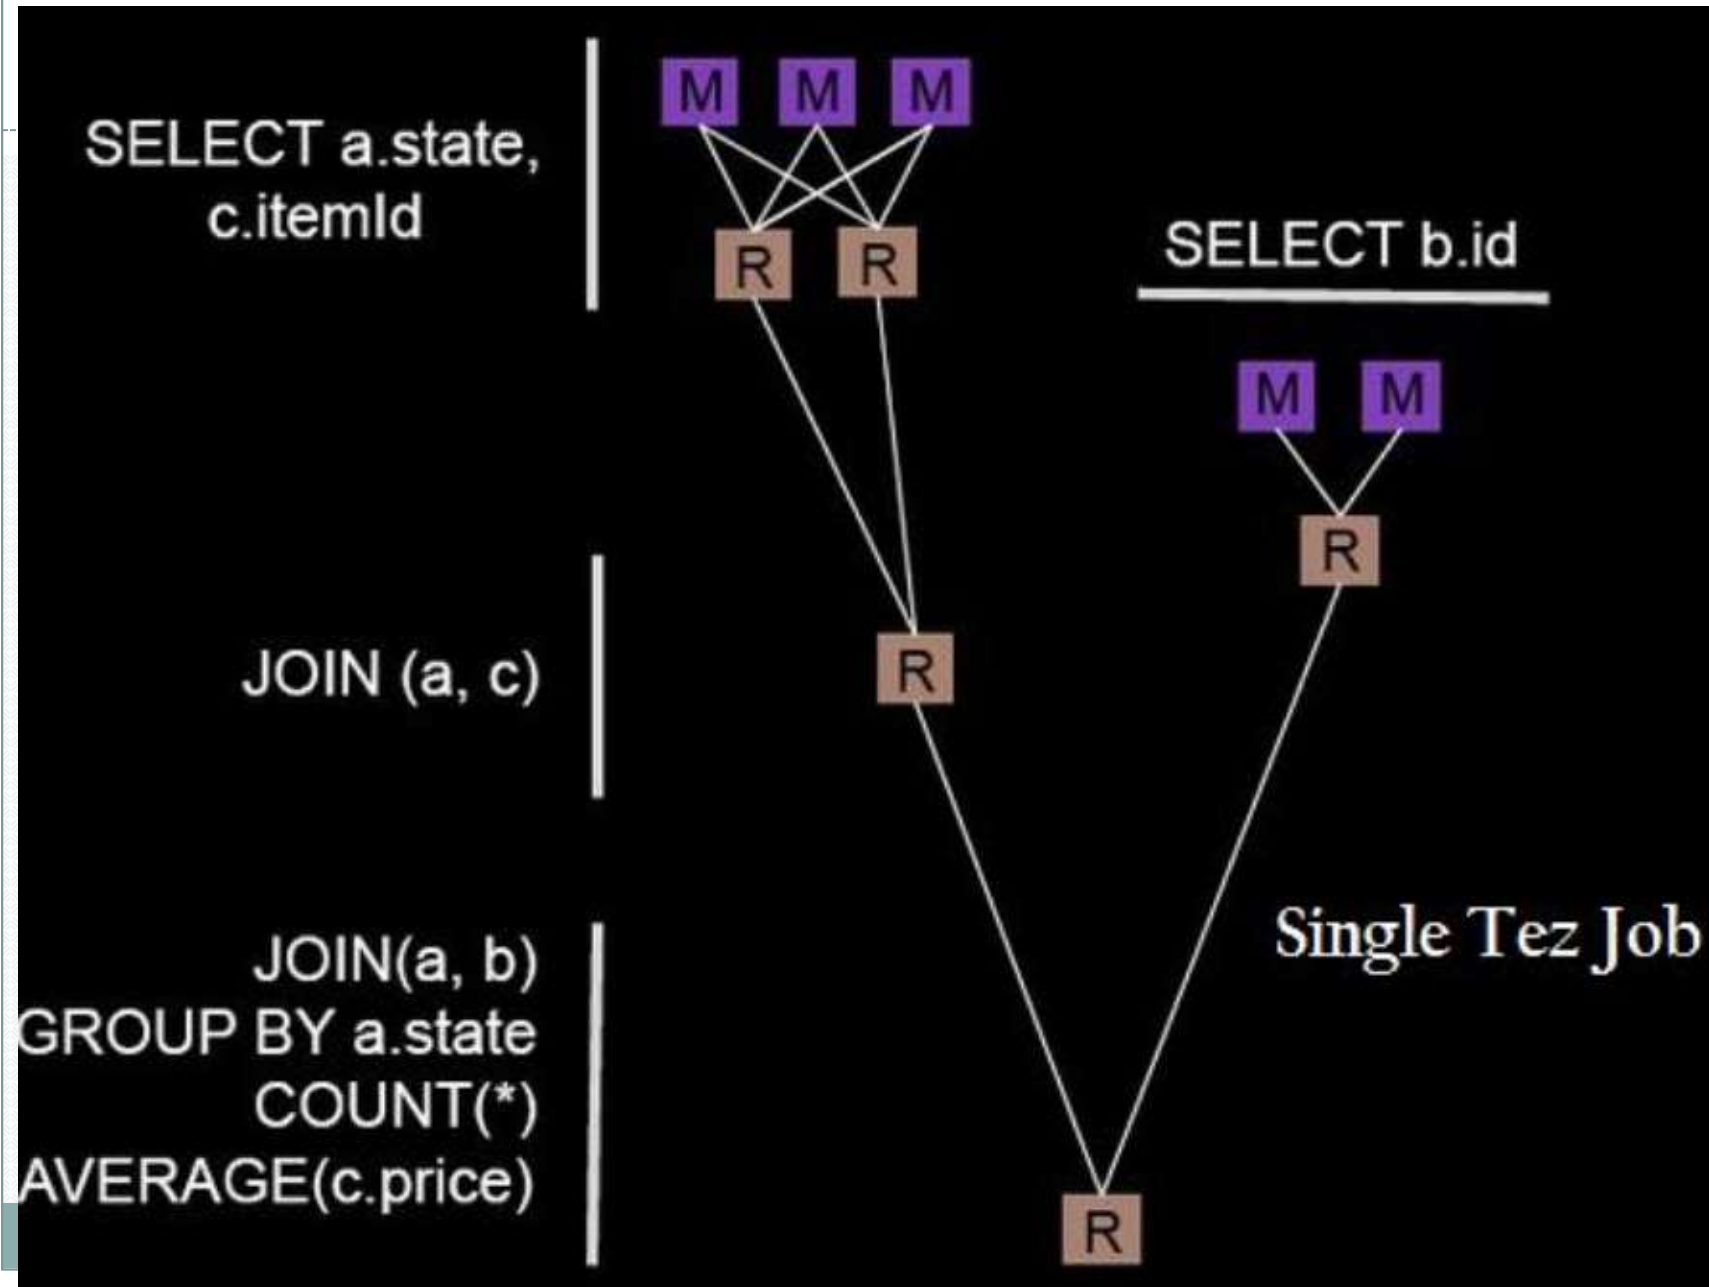
```
SELECT a.state, COUNT(*), AVERAGE(c.price)
FROM a
JOIN b ON(a.id = b.id)
JOIN c ON(a.itemId = c.itemId)
GROUP BY a.state
```



# Query in MapReduce

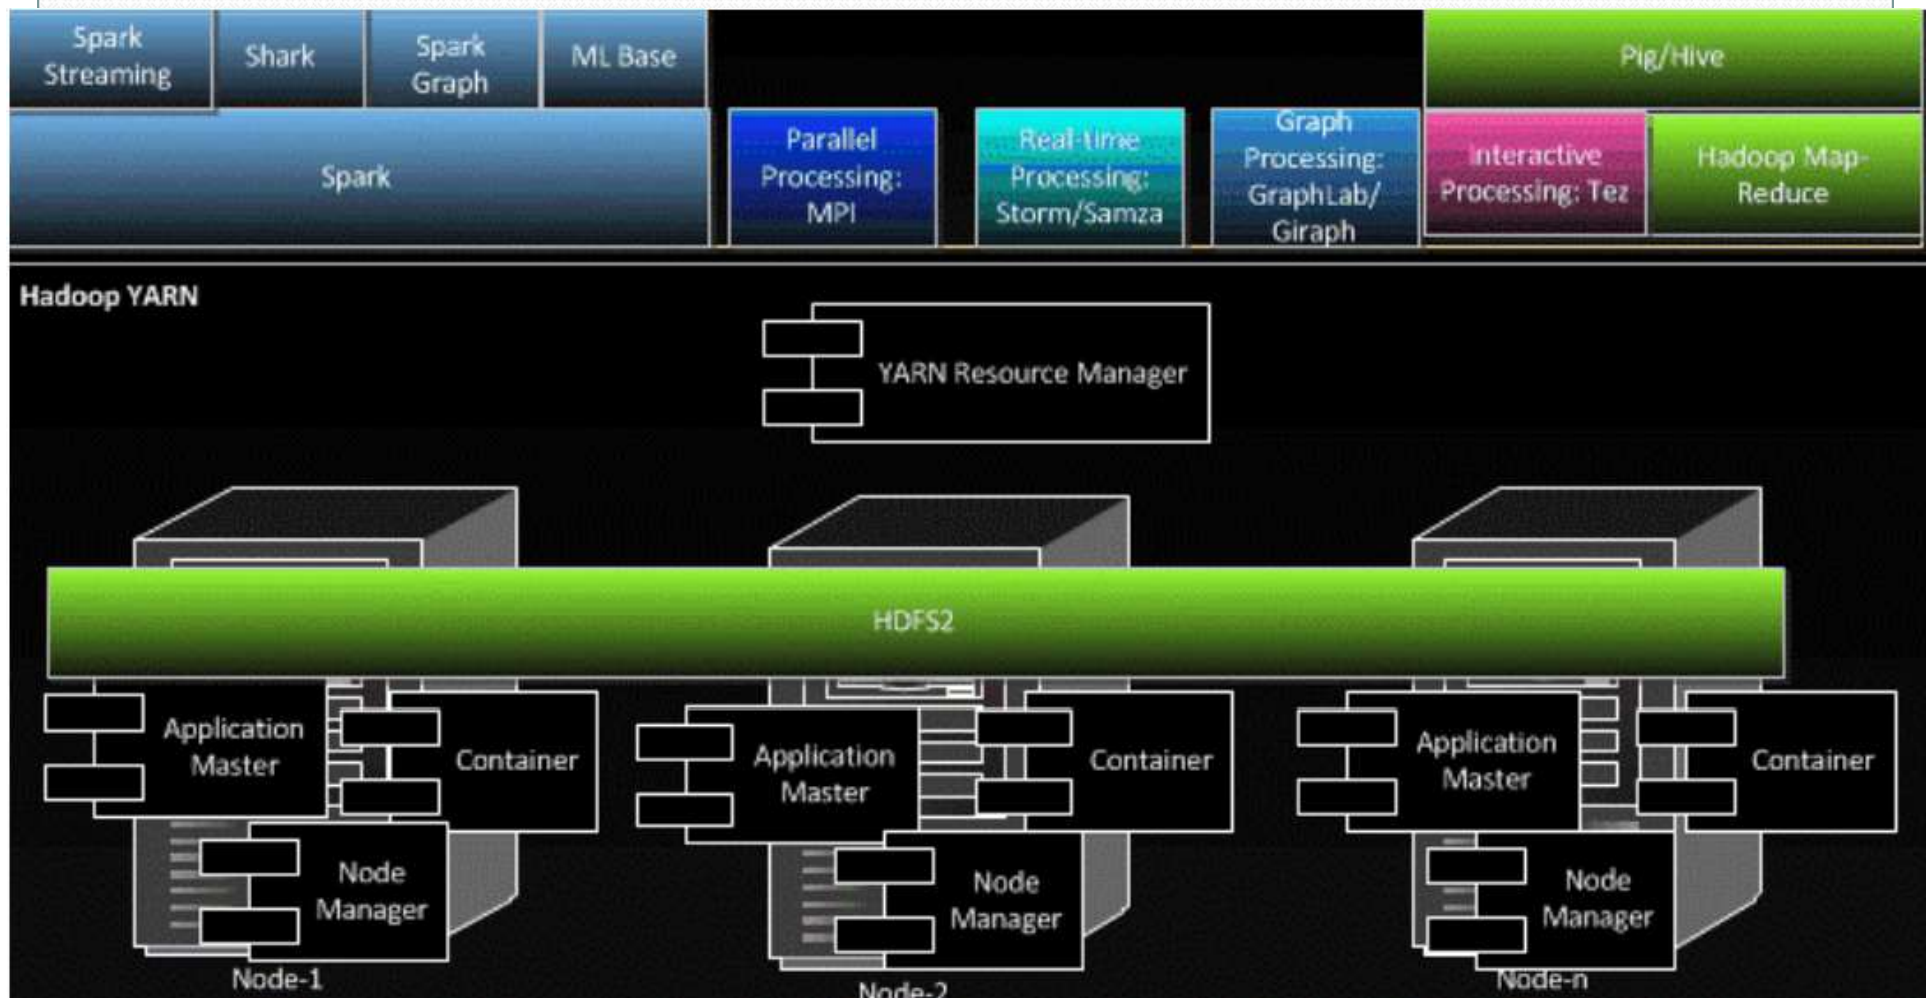


# Query in Apache Tez





# Expanded YARN Framework



# Tez Summary



- Tez development framework includes Java APIs
  - Provide developers with an intuitive interface to create the data-processing graphs for the applications' data-processing flows.
- Tez provides additional APIs to inject custom logic
- Modular input, output, and processor elements are provided.
- Applications developed via these APIs execute efficiently in a Hadoop environment
- Tez framework provides the interaction with other stack components
- Tez applications represent customized, optimized, and YARN integrated solutions that take advantage of Hadoop's scalable and fault-tolerant runtime environment.



# Tez Summary



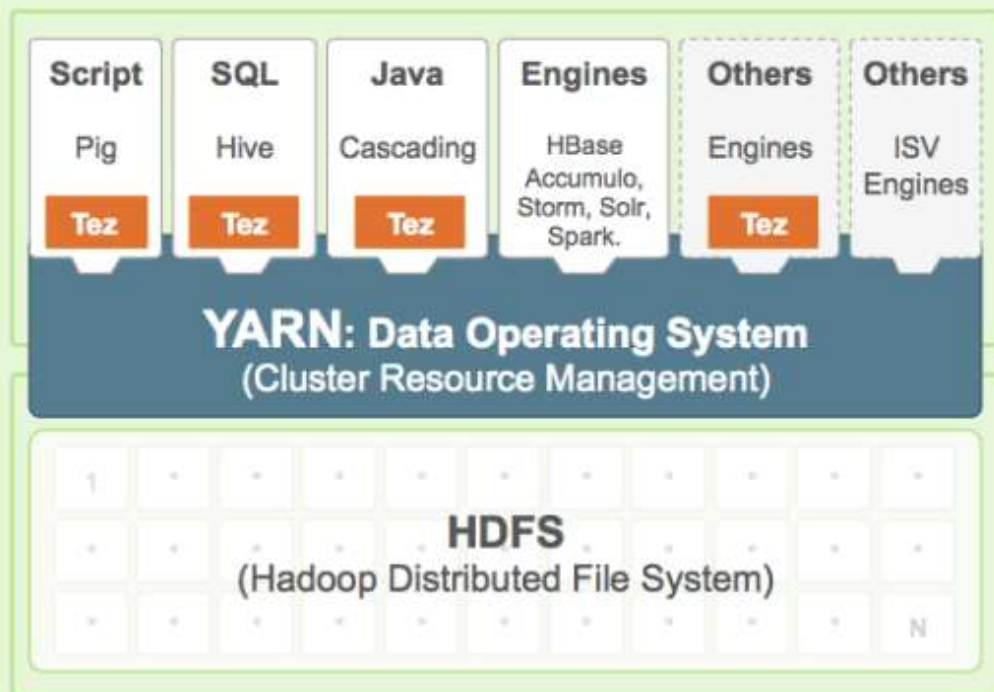
## Hadoop 1

- Silos & Largely batch
- Single Processing engine



## Hadoop 2 w/ Tez

- Multiple Engines, Single Data Set
- Batch, Interactive & Real-Time





# Class 4

# Apache Pig

PLATFORM FOR  
ANALYZING LARGE  
DATA SETS



# Summary



- Big demand for parallel data processing
  - Emerging tools that do not look like SQL DBMS
  - Programmers like dataflow pipes over static files
- Hence the excitement about Map-Reduce
- But, Map-Reduce is too low-level and rigid

**Pig Latin**

Sweet spot between map-reduce and SQL

# What is Pig

14



- A platform for analyzing large data sets
  - Contains a high-level language for **expressing** data analysis programs.
- Compiles down to MapReduce jobs
- Developed by Yahoo!
- Open-source language



# Pig Components

15

## Two Main Components

- High-level language (Pig Latin)
  - Set of commands
- Two execution modes
  - Local: reads/write to local file system
  - Mapreduce: connects to Hadoop cluster and reads/writes to HDFS

## Two modes

- Interactive mode
  - Console
- Batch mode
  - Submit a script

# Pig's Team



In 2009, 8 out of 10 Pig contributors were from Yahoo.



# Apache Pig Background



- Yahoo! was the first big adopter of Hadoop.
- Hadoop gained popularity in the company quickly.
- Yahoo! Research developed Pig to address the need for a higher level language.
- Roughly 30% of Hadoop jobs run at Yahoo! are Pig jobs.



# How Apache Pig is used?



- Web log processing
- Data processing for web search platforms
- Ad hoc queries for large datasets
- Rapid prototyping of algorithms for large datasets



# Apache Pig

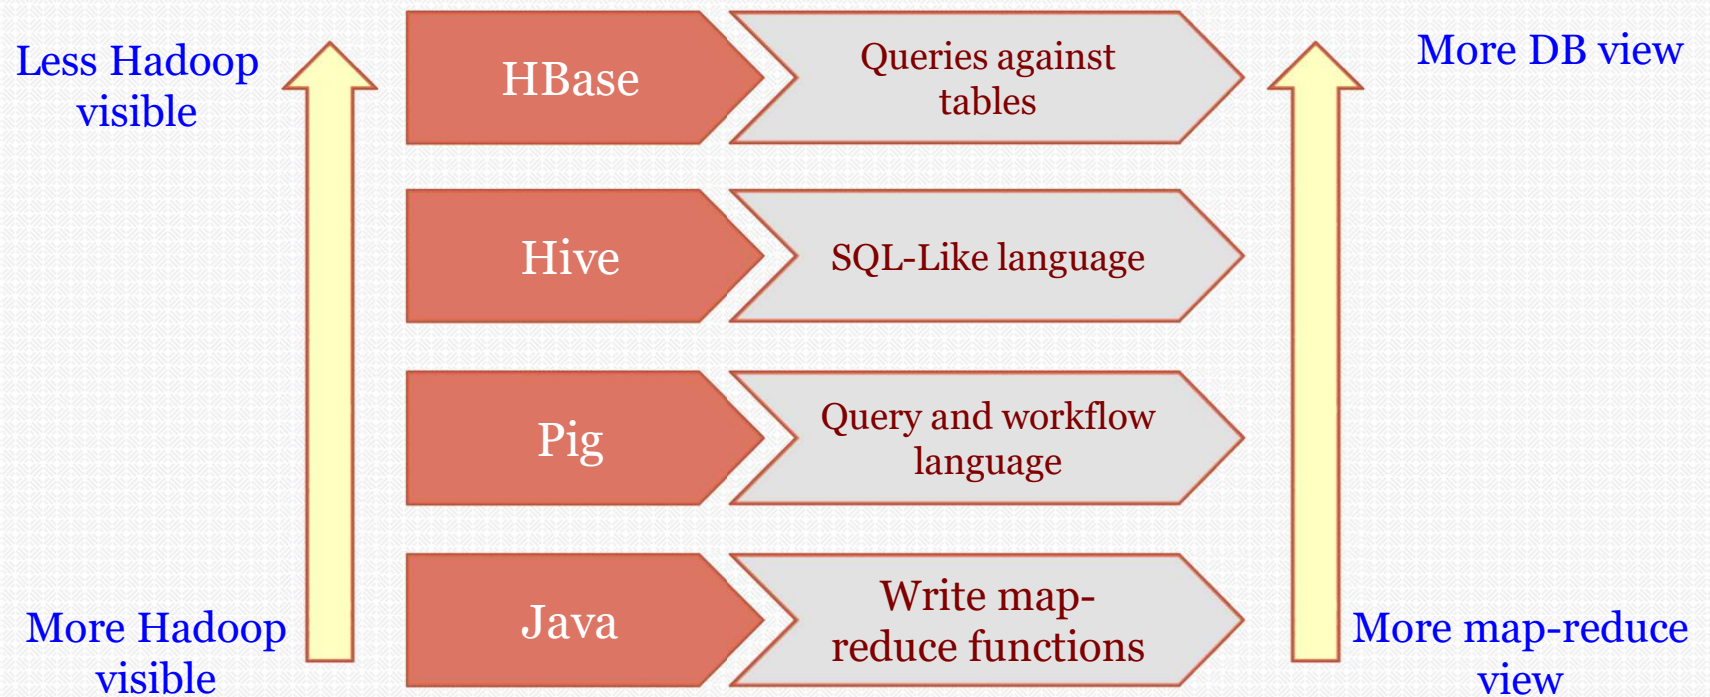


- High level scripting language in Apache Hadoop stack
- Good for data analysis as data flows.
- Has many data manipulation commands
- User Defined Functions(UDF)
  - invoke code in many languages
  - JRuby, Jython and Java.
- Can run Pig scripts from other languages.
- Pig may be used as a component for large and complex applications



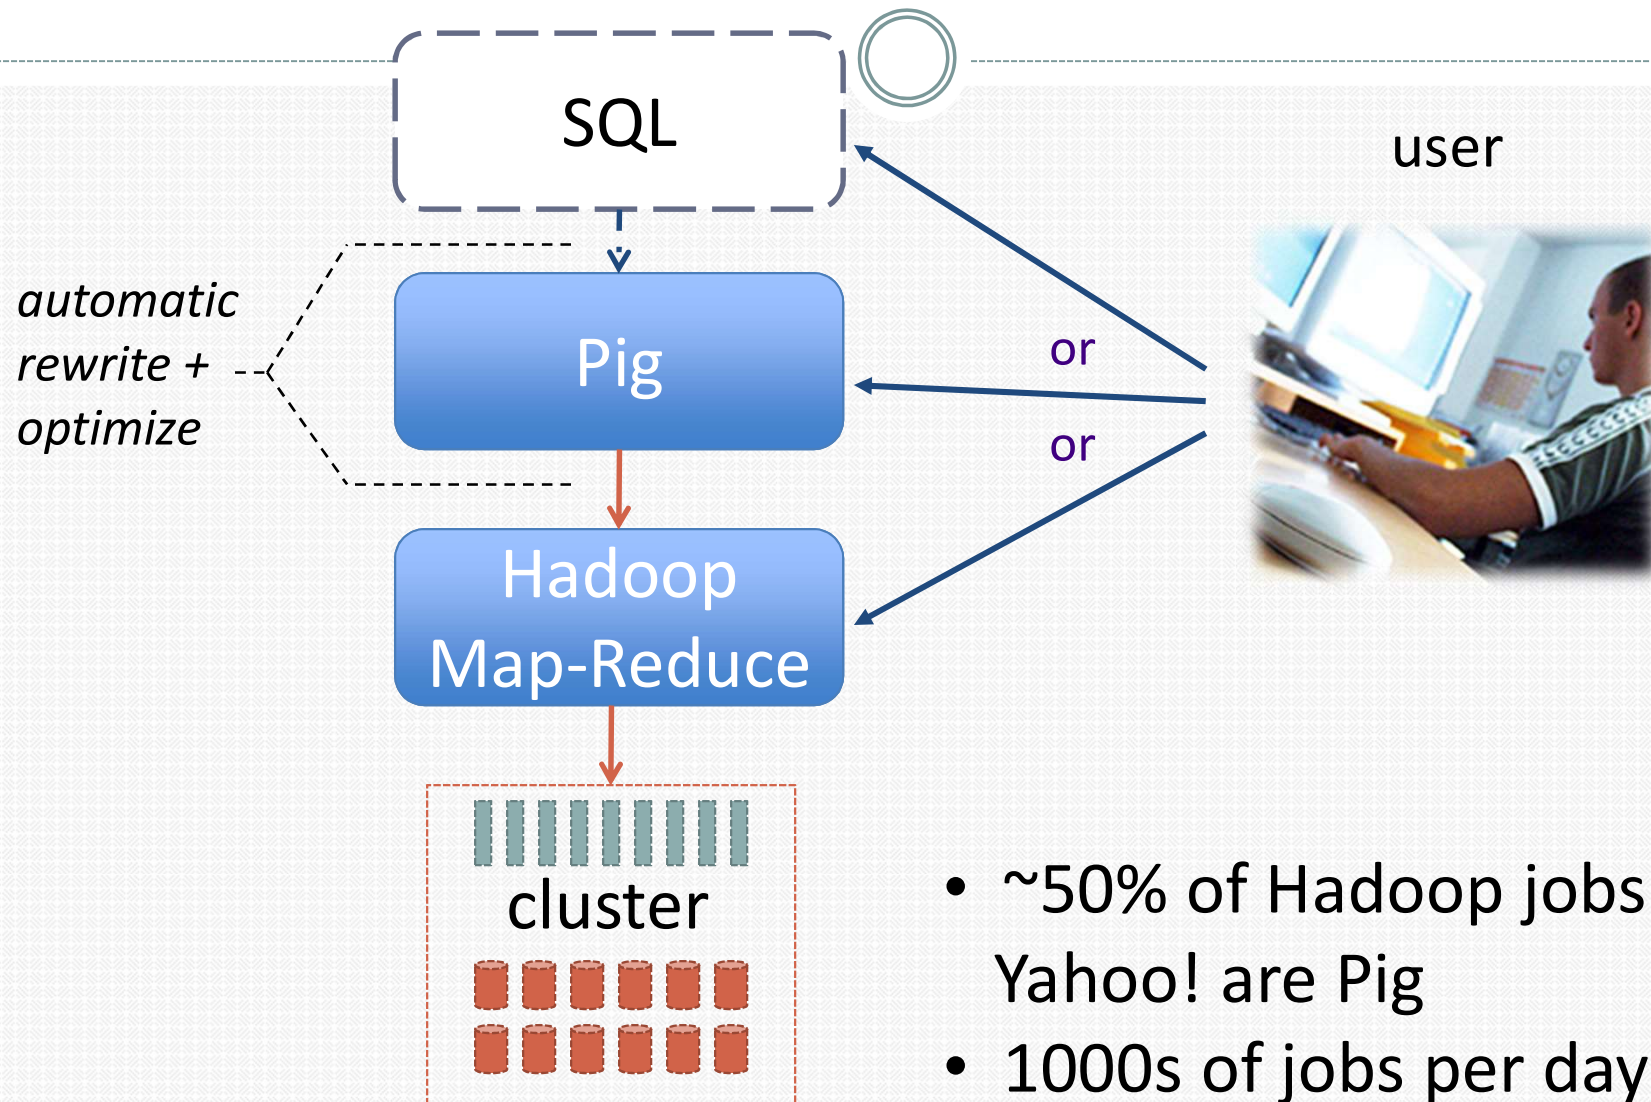
# Levels of Abstraction

20





# Implementation



- ~50% of Hadoop jobs at Yahoo! are Pig
- 1000s of jobs per day

# Pig Latin vs. SQL

22

- Pig Latin is procedural (dataflow programming model)
  - Step-by-step query style is much cleaner and easier to write
- SQL is declarative but not step-by-step style

SQL

```
insert into ValuableClicksPerDMA
select dma, count(*)
from geoinfo join (
    select name, ipaddr
    from users join clicks on (users.name = clicks.user)
    where value > 0;
) using ipaddr
group by dma;
```

Pig  
Latin

```
Users          = load 'users' as (name, age, ipaddr);
Clicks         = load 'clicks' as (user, url, value);
ValuableClicks = filter Clicks by value > 0;
UserClicks     = join Users by name, ValuableClicks by user;
Geoinfo        = load 'geoinfo' as (ipaddr, dma);
UserGeo        = join UserClicks by ipaddr, Geoinfo by ipaddr;
ByDMA          = group UserGeo by dma;
ValuableClicksPerDMA = foreach ByDMA generate group, COUNT(UserGeo);
store ValuableClicksPerDMA into 'ValuableClicksPerDMA';
```



# Pig Latin vs. SQL

23

- **In Pig Latin**

- Lazy evaluation (data not processed prior to STORE command)
- Data can be stored at any point during the pipeline
- Schema and data types are lazily defined at run-time
- An execution plan can be explicitly defined
  - Use optimizer hints
  - Due to the lack of complex optimizers

- **In SQL:**

- Query plans are solely decided by the system
- Data cannot be stored in the middle
- Schema and data types are defined at the creation time





# Apache Pig




- Map Reduce is very powerful, but:
  - Requires a Java programmer.
  - User has to write common SQL functionality from scratch
    - ✦ joins, filtering, etc.
- Pig provides a higher level language, Pig Latin:
  - Increases productivity.
  - 10 lines of Pig Latin  $\approx$  200 lines of Java.
  - 4 hours in Java  $\approx$  15 minutes in Pig Latin.
  - Good for non-Java programmers.
  - Provides common operations:
    - ✦ Join, Group, Filter, Sort.

# Java vs. Pig

25


**1/20 the lines of code**



Java

Pig

**1/16 the development time**



Java

Pig

Performance is comparable (Java is slightly better)





# Pig Lating



- Pig Latin is a data flow language
  - Neither procedural nor declarative.
- User code and existing binaries can be included almost anywhere.
- Metadata not required, but used when available.
- Supports nested types.
- Operates on files in HDFS
  - No need for table mappings
  - No need for ETL



# Several ways to use Pig



- Type commands in **Pig's Grunt** (Pig CLI)
- Submit a script to Pig
- Use it from Java with PigServer Java class and JDBC-like interface
- Use it from Eclipse with PigPen
  - PigPen is an Eclipse plugin
  - Allows textual and graphical scripting
  - Samples data and shows example data flow

# Grunt is a Pig's CLI



```
-bash-4.1$ pig
WARNING: Use "yarn jar" to launch YARN applications.
16/02/15 23:04:02 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
16/02/15 23:04:02 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
16/02/15 23:04:02 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
2016-02-15 23:04:02,623 [main] INFO  org.apache.pig.Main - Apache Pig version 0.
15.0.2.3.4.0-3485 (rexported) compiled Dec 16 2015, 04:30:33
2016-02-15 23:04:02,623 [main] INFO  org.apache.pig.Main - Logging error message
s to: /home/st65/pig_1455595442621.log
2016-02-15 23:04:02,651 [main] INFO  org.apache.pig.impl.util.Utils - Default bo
otup file /home/st65/.pigbootup not found
2016-02-15 23:04:03,106 [main] INFO  org.apache.pig.backend.hadoop.executionengi
ne.HExecutionEngine - Connecting to hadoop file system at: hdfs://hphone:8020
2016-02-15 23:04:03,943 [main] INFO  org.apache.pig.PigServer - Pig Script ID fo
r the session: PIG-default-5f3f236a-b708-45ce-b025-1a82c83b942f
2016-02-15 23:04:04,542 [main] INFO  org.apache.hadoop.yarn.client.api.impl.Time
lineClientImpl - Timeline service address: http://hpthree:8188/ws/v1/timeline/
2016-02-15 23:04:04,771 [main] INFO  org.apache.pig.backend.hadoop.ATSService -
Created ATS Hook
grunt> █
```



# How Apache Pig works



## Pig Latin

```
A = LOAD 'myfile'
  AS (x, y, z);
B = FILTER A by x > 0;
C = GROUP B BY x;
D = FOREACH A GENERATE
  x, COUNT(B);
STORE D INTO 'output';
```



pig.jar:

- parses
- checks
- optimizes
- plans execution
- submits jar to Hadoop
- monitors job progress

Execution Plan  
Map:  
Filter

Reduce:  
Count





# Example. Pig reads and displays a file



```
grunt> log = load 'excite-small.log' as (user, timestamp, query);  
grunt> dump log
```

- Pig reads the **excite-small.log** file from HDFS
  - It takes it from your HDFS user directory
- Pig translates the command into MapReduce jobs
  - The jobs are executed in parallel in the cluster
- Output looks like this:

(818D7157855D5D48,970916080224,halloween )

(99BA461C4F96233F,970916023950,"tibetan terrier")

(887E337AB02C2BF7,970916170230,edgar allen poe)



# Pig Data Types



- **Scalar types:**
  - Int
  - Long
  - Double
  - Chararray
  - Bytearray
- **Complex types:**
  - Map: associative array
  - Tuple: ordered list of data, elements may be of any scalar or complex type
  - Bag: unordered collection of tuples



## Pig Example (continued)



```
log = load 'excite-small.log' as (user, timestamp, query);  
grp = group log by user;  
dump grp;
```

--- OUTPUT: -----

```
(74BB49A63CoC4D79,  
{(74BB49A63CoC4D79,970916104542,mcgraw hill),  
(74BB49A63CoC4D79,970916073222,green tree),  
(74BB49A63CoC4D79,970916073158,green tree),  
(74BB49A63CoC4D79,970916071727,be-cubed),  
(74BB49A63CoC4D79,970916071706,be-cubed),  
(74BB49A63CoC4D79,970916104549,mcgraw hill),  
(74BB49A63CoC4D79,970916071651,b-cubed)}  
)
```



# Apache Pig Example



- Aggregation

Let's count the number of times each user appears in the excite data set.

```
log = LOAD 'excite-small.log'
      AS (user, timestamp, query);
grpd = GROUP log BY user;
cntd = FOREACH grpd GENERATE group, COUNT(log);
STORE cntd INTO 'output';
```

## Results:

002BB5A52580A8ED	18
005BD9CD3AC6BB38	18

## Another Example in Pig Latin



```
visits      = load '/data/visits' as (user, url, time);
gVisits     = group visits by url;
visitCounts = foreach gVisits generate url, count(visits);

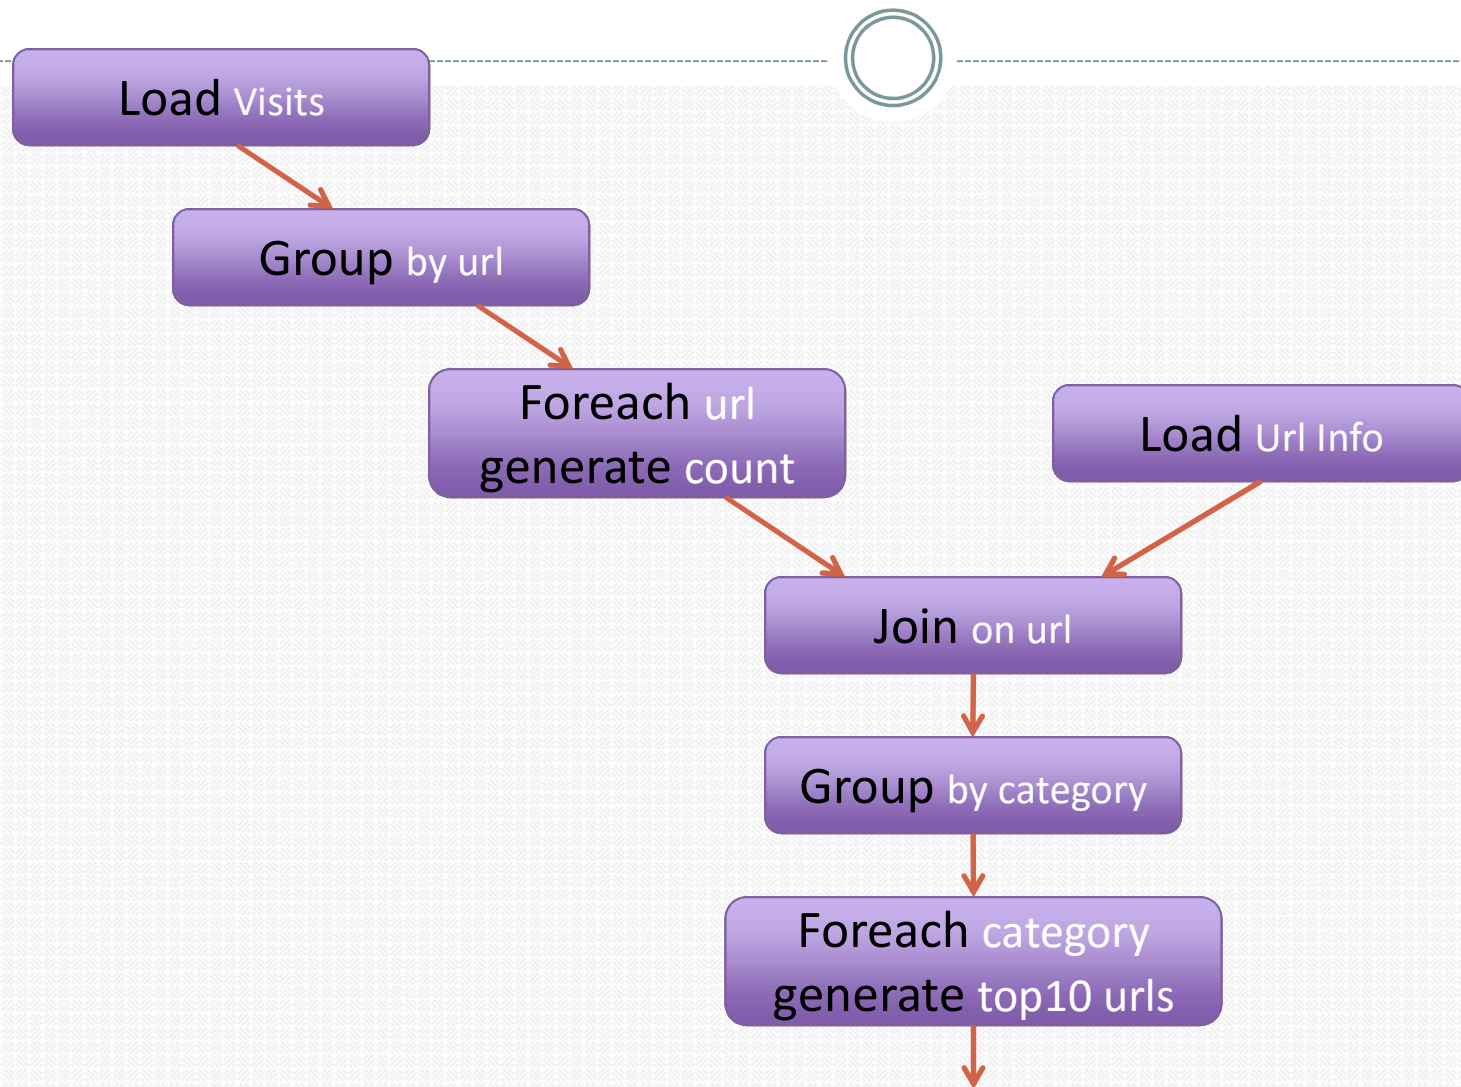
urlInfo     = load '/data/urlInfo' as (url, category, pRank);
visitCounts = join visitCounts by url, urlInfo by url;

gCategories = group visitCounts by category;
topUrls     = foreach gCategories generate top(visitCounts,10);

store topUrls into '/data/topUrls';
```

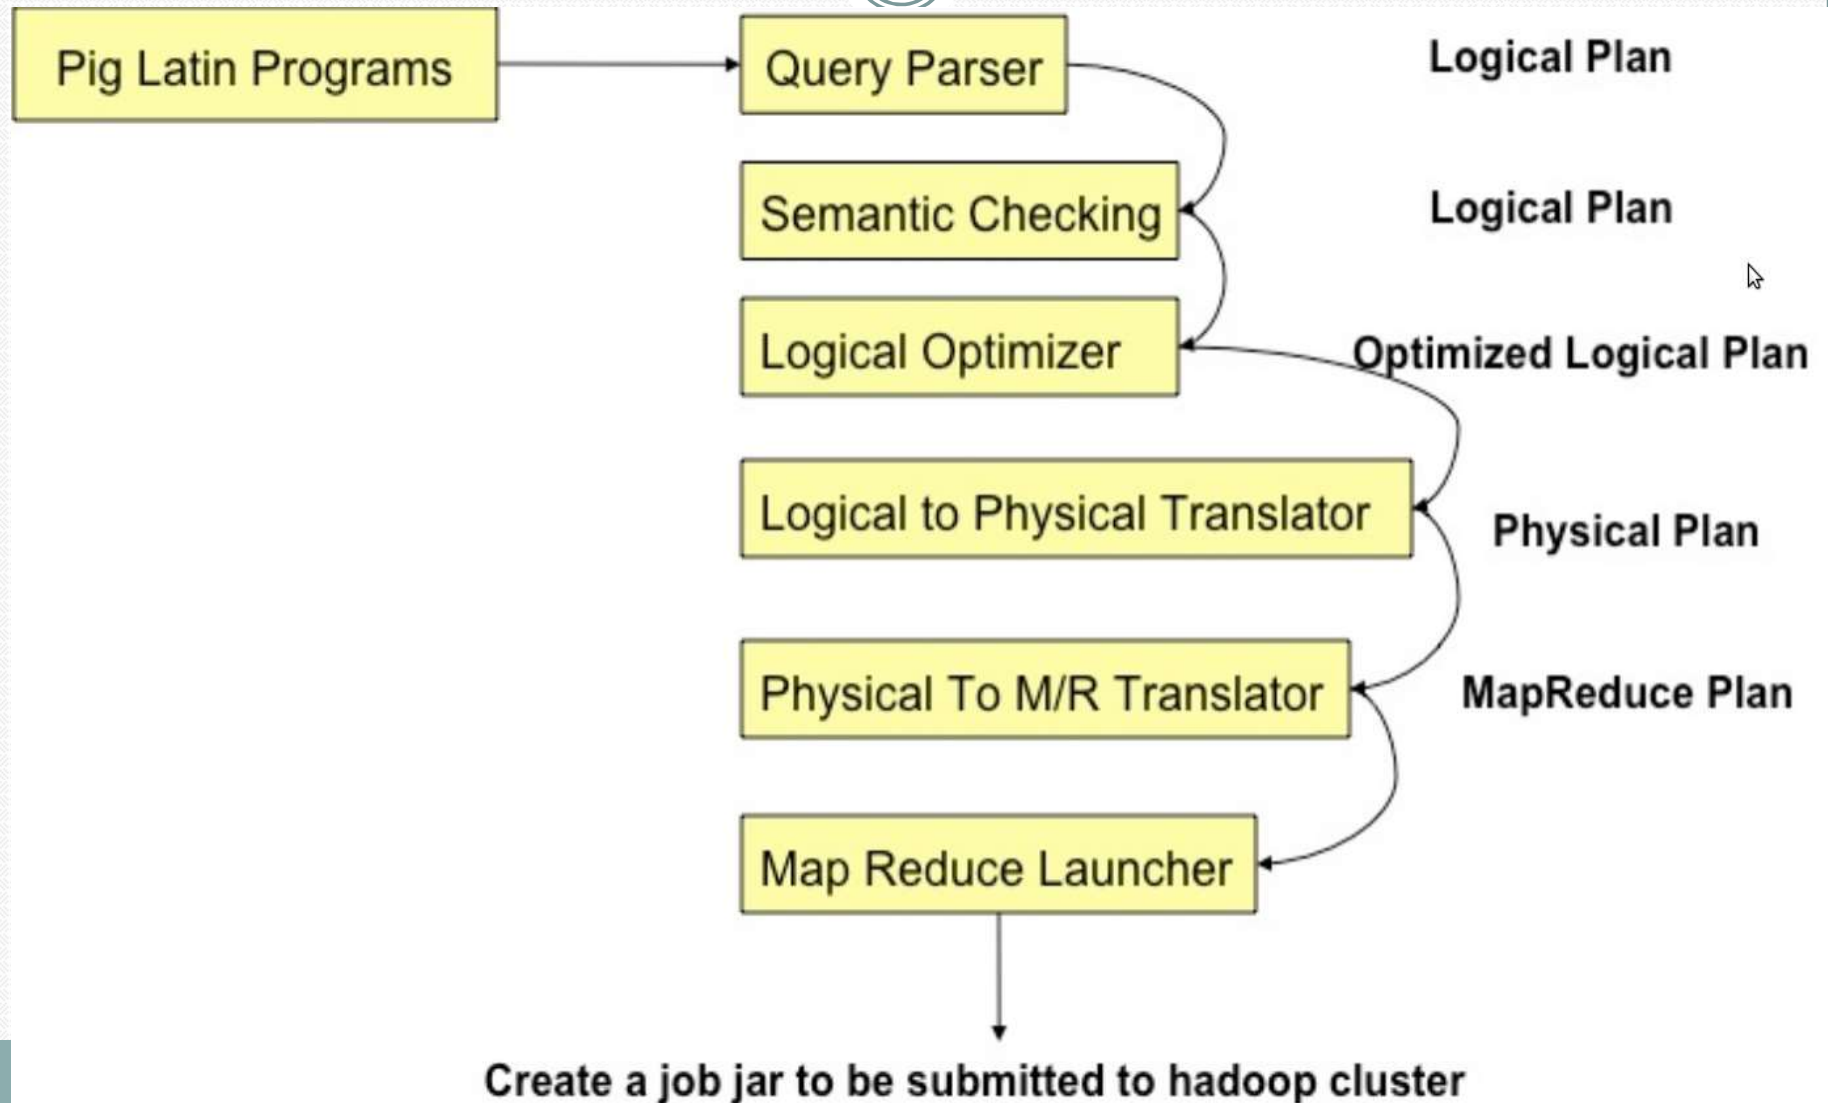


# This is how it is executed



# Pig Compilation


36





# Logic Plan

```
A=LOAD 'file1' AS (x, y, z);  
B=LOAD 'file2' AS (t, u, v);  
C=FILTER A by y > 0;  
D=JOIN C BY x, B BY u;  
E=GROUP D BY z;  
F=FOREACH E GENERATE  
    group, COUNT(D);  
STORE F INTO 'output';
```



# Physical Plan

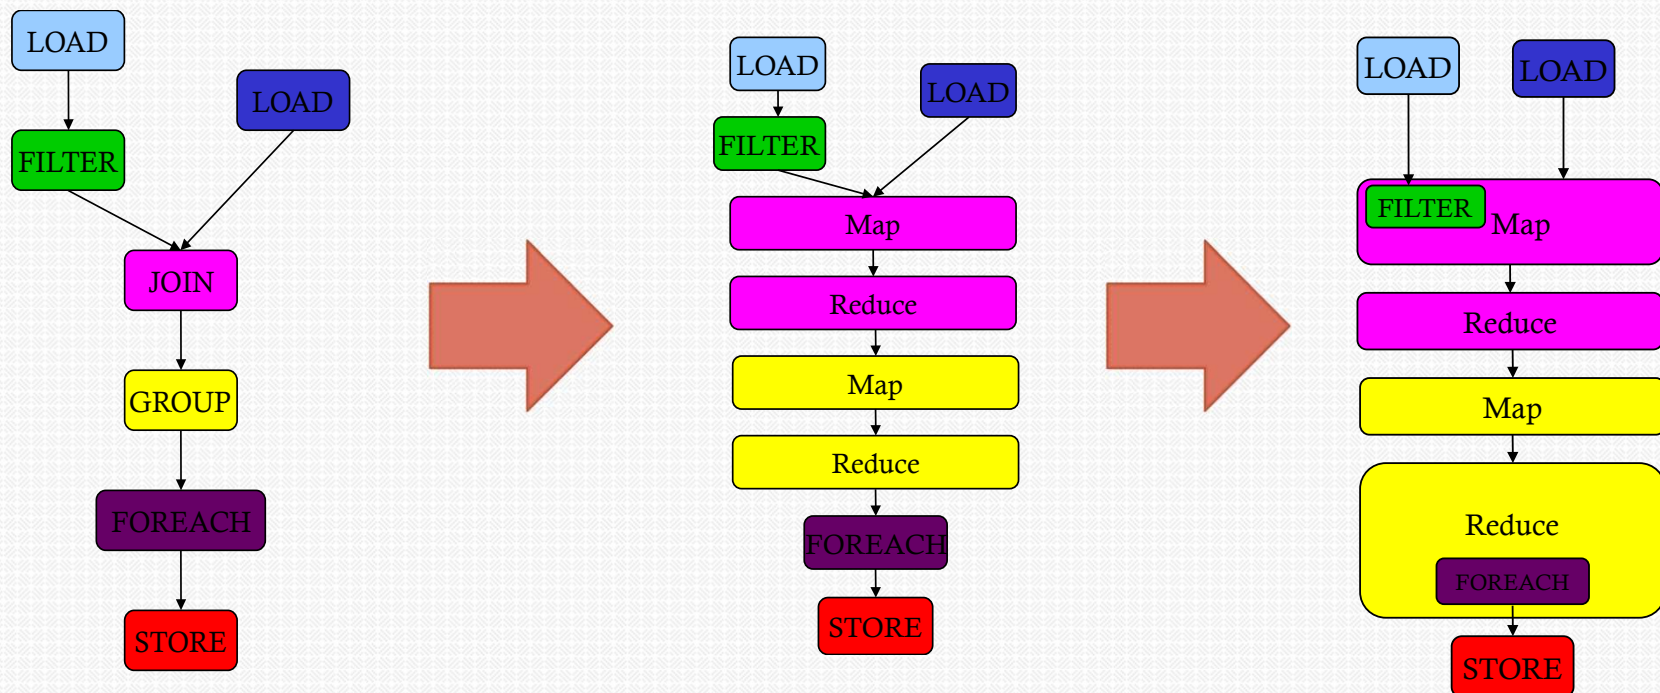
38

- 1:1 correspondence with the logical plan
- **Except for:**
  - Join, Distinct, (Co)Group, Order
- Several optimizations are done automatically



# Generation of Physical Plans

39



If the Join and Group By are on the same key → The two map-reduce jobs would be merged into one.



# - Grouping

- In Pig grouping is separate operation from applying aggregate functions.
- In previous slide, output of group statement is (key, bag), where key is the group key and bag contains a tuple for every record with that key.
- For example:

alan 1

bob 9       =>     alan, { (alan, 1), (alan, 3) }

alan 3             bob, { (bob, 9) }





# Filtering

Now let's apply a filter to the groups so that we only get the high frequency users.

```
log      = LOAD 'excite-small.log'
          AS (user, time, query);
grpd     = GROUP log BY user;
cntd     = FOREACH grpd GENERATE
          group, COUNT(log) AS cnt;
fltrd    = FILTER cntd BY cnt > 50;
STORE fltrd INTO 'output';
```

## Results:

0B294E3062F036C3	61
128315306CE647F6	78
7D286B5592D83BBE	59



# Ordering

Sort the high frequency users by frequency.

```
log    = LOAD 'excite-small.log'
        AS (user, time, query);
grpd   = GROUP log BY user;
cntd   = FOREACH grpd GENERATE
        group, COUNT(log) AS cnt;
fltrd  = FILTER cntd BY cnt > 50;
srtd   = ORDER fltrd BY cnt;
STORE srtd INTO 'output';
```

## Results:

7D286B5592D83BBE	59
0B294E3062F036C3	61
128315306CE647F6	78





# Prepare Data For Joins

Run a pig job to prepare the Shakespeare and King James Bible data for the following examples:

```
A = load 'bible';  
B = foreach A generate  
    flatten(TOKENIZE((chararray)$0)) as word;  
C = filter B by word matches '\\w+';  
D = group C by word;  
E = foreach D generate COUNT(C), group;  
store E into 'bible_freq';
```

Repeat for Shakespeare, changing 'bible' to 'input' in line 1 and 'bible\_freq' to 'shakespeare\_freq' in line 7.



# Join

We can use join to find words that are in both the King James version of the Bible and Shakespeare's collected works.

```
bard    = LOAD 'shakespeare_freq' AS (freq, word);  
kjv     = LOAD 'bible_freq' AS (freq, word);  
inboth  = JOIN bard BY word, kjv BY word;  
STORE inboth INTO 'output';
```

## Results:

2	Abide	1	Abide
2	Abraham	111	Abraham
3	...		





# Anti-Join

Find words that are in the Bible that are not in Shakespeare.

```
bard    = LOAD 'shakespeare_freq' AS (freq, word);  
kjb     = LOAD 'bible_freq' AS (freq, word);  
grpd    = COGROUP bard BY word, kjb BY word;  
nobard  = FILTER grpd BY COUNT(bard) == 0;  
out     = FOREACH nobard GENERATE FLATTEN(kjb);  
STORE out INTO 'output';
```

## Results:

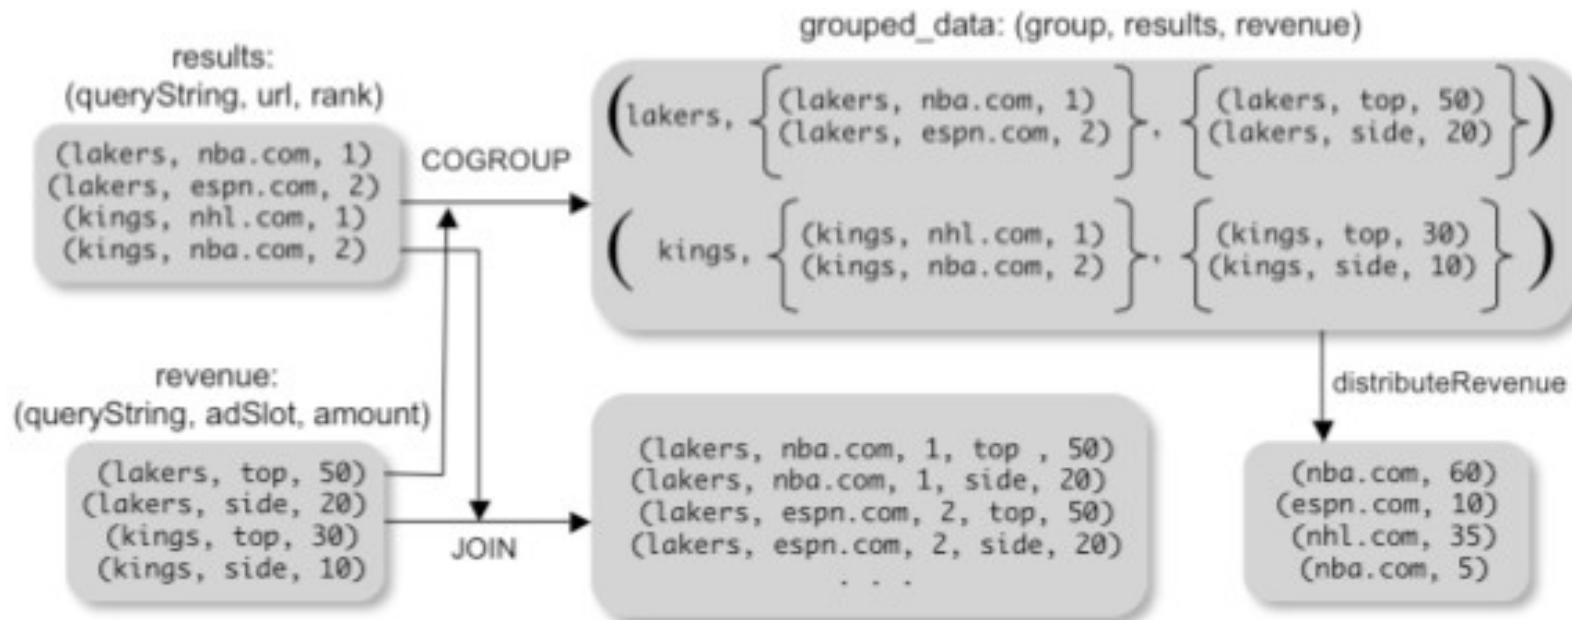
1	Abimael
22	Abimelech

# CoGroup: combination of **join** and **group by**

```
grouped_data = COGROUP results BY queryString,  
                    revenue BY queryString;
```

```
join_result = JOIN results BY queryString,  
                    revenue BY queryString;
```

```
temp_var = COGROUP results BY queryString,  
                    revenue BY queryString;  
join_result = FOREACH temp_var GENERATE  
                    FLATTEN(results), FLATTEN(revenue);
```







# Cogrouping

- Cogrouping is a generalization of grouping.
- Keys for two (or more) inputs are collected.
- In previous slide, output of group statement is (key, bag1, bag2), where key is the group key, bag1 contains a tuple for every record in the first input with that key, and bag2 contains a tuple for every record in the second input with that key.



# Cogrouping Continued

- For example:

Input 1:

alan 1

bob 9

alan 3

Input 2:

alan 5

bob 9

Output:

alan, {(alan, 1), (alan, 3)}, {(alan, 5)}

bob, {(bob, 9)}, {}

john, {}, {(john, 3)}





# Using Types

- By default Pig treats data as un-typed.
- User can declare types of data at load time.

```
log = LOAD 'shakespeare_freq'  
      AS (freq:int, word: chararray);
```

- If data type is not declared but script treats value as a certain type, Pig will assume it is of that type and cast it.

```
log      = LOAD 'shakespeare_freq' AS (freq, word);  
weighted = FOREACH log  
            GENERATE freq * 100; --freq cast to int
```



# Custom Load & Store

- By default Pig assumes data is tab separated UTF-8.
- User can set a different delimiter.
- If you have a different format, you use another load/store function to handle de/serialization of the data

```
A = LOAD 'data.json'  
    USING PigJsonLoader();
```

...

```
STORE INTO 'output.json'  
    USING PigJsonLoader();
```





# User Defined Functions

- For logic that cannot be done in Pig.
- Can be used to do column transformation, filtering, ordering, custom aggregation.
- For example, you want to write custom logic to do user session analysis:

```
log  = LOAD 'excite-small.log'
      AS (user, time, query);
grpd = GROUP log BY user;
cntd = FOREACH grpd GENERATE
      group, SessionAnalysis(log);
STORE cntd INTO 'output';
```



# Nested Operations

- A FOREACH can apply a set of operators to every tuple in turn.
- Using the previous session analysis example, assume the UDF requires input to be sorted by time:

```
log  = LOAD 'excite-small.log'
      AS (user, time, query);
grpd = GROUP log BY user;
cntd = FOREACH grpd {
      srtd = ORDER log BY time;
      GENERATE group, SessionAnalysis(srtd);
}
STORE cntd INTO 'output';
```





# Split

- Data flow need not be linear:

```
A = LOAD 'data';  
B = GROUP A BY $0;  
...  
C = GROUP A by $1;  
...
```

- Split can be used explicitly:

```
A = LOAD 'data';  
SPLIT A INTO NEG IF $0 < 0, POS IF $0 > 0;  
• B = FOREACH NEG GENERATE ...  
...
```



# Pig Commands

Pig Command	What it does
load	Read data from file system.
store	Write data to file system.
foreach	Apply expression to each record and output one or more records.
filter	Apply predicate and remove records that do not return true.
group/cogroup	Collect records with the same key from one or more inputs.
join	Join two or more inputs based on a key.
order	Sort records based on a key.
distinct	Remove duplicate records.
union	Merge two data sets.
split	Split data into 2 or more sets, based on filter conditions.
stream	Send all records through a user provided binary.
dump	Write output to stdout.
limit	Limit the number of records.



# Example I: More Details

55

Read file from HDFS

The input format (text, tab delimited)

Define run-time schema

```
raw = LOAD 'excite.log' USING PigStorage('\t') AS (user, id, time, query);  
clean1 = FILTER raw BY id > 20 AND id < 100;  
clean2 = FOREACH clean1 GENERATE  
        user, time,  
        org.apache.pig.tutorial.sanitize(query) as query;  
user_groups = GROUP clean2 BY (user, query);  
user_query_counts = FOREACH user_groups GENERATE group, COUNT(clean2), MIN(clean2.time), MAX(clean2.time);  
STORE user_query_counts INTO 'uq_counts.csv' USING PigStorage(',');
```

Filter the rows on predicates

For each row, do some transformation

Grouping of records

Compute aggregation for each group

Store the output in a file

Text, Comma delimited

# Pig: Language Features

56

- **Keywords**
  - Load, Filter, Foreach Generate, Group By, Store, Join, Distinct, Order By, ...
- **Aggregations**
  - Count, Avg, Sum, Max, Min
- **Schema**
  - Defines at query-time not when files are loaded
- UDFs
- Packages for common input/output formats



# Example 2

57

Script can take arguments

Data are “ctrl-A” delimited

Define types of the columns

```
A = load '$widerow' using PigStorage('\u0001') as (name: chararray, c0: int, c1: int, c2: int);  
B = group A by name parallel 10;  
C = foreach B generate group, SUM(A.c0), SUM(A.c1), SUM(A.c2) as c2;  
D = filter C by c0 > 100 and c1 > 100 and c2 > 100;  
store D into '$out';
```

Specify the need of 10 reduce tasks

# Example 3: Re-partition Join

58

Register UDFs & custom inputformats

```
register pigperf.jar;  
A = load 'page_views' using  
org.apache.pig.test.udf.storefunc.PigPerformanceLoader()  
    as (user, action, timespent, query_term timestamp,  
estimated_revenue);  
B = foreach A generate user, (double) estimated_revenue;  
alpha = load 'users' using PigStorage('\u0001') as (name,  
phone, address, city, state, zip);  
beta = foreach alpha generate name, city;  
C = join beta by name, B by user parallel 40;  
D = group C by $0;  
E = foreach D generate group, SUM(C.estimated_revenue);  
store E into 'L3out';
```

Function the jar file to read the input file

Load the second file

Join the two datasets (40 reducers)

Group after the join (can reference columns by position)

This grouping can be done in the same map-reduce job because it is on the same key (Pig can do this optimization)

# Example 4: Replicated Join

59

```
register pigperf.jar;  
A = load 'page_views' using  
org.apache.pig.test.udf.storefunc.PigPerformanceLoader()  
    as (user, action, timespent, query_term, timestamp,  
estimated_revenue);  
Big = foreach A generate user, (double) estimated_revenue;  
alpha = load 'users' using PigStorage('\u0001') as (name,  
phone, address, city, state, zip);  
small = foreach alpha generate name, city;  
C = join Big by user, small by name using 'replicated';  
store C into 'out';
```

Map-only join (the small dataset is the second)

Optimization in joining a big dataset with a small  
one



# Example 5: Multiple Outputs

60

```
A = LOAD 'data' AS (f1:int,f2:int,f3:int);
```

```
DUMP A;
```

```
(1,2,3)
```

```
(4,5,6)
```

```
(7,8,9)
```

Split the records into sets

```
SPLIT A INTO X IF f1<7, Y IF f2==5, Z IF (f3<6 OR f3>6);
```

```
DUMP X;
```

```
(1,2,3)
```

```
(4,5,6)
```

Dump command to display the data

```
DUMP Y;
```

```
(4,5,6)
```

Store multiple outputs

```
STORE x INTO 'x_out';
```

```
STORE y INTO 'y_out';
```

```
STORE z INTO 'z_out';
```

# Run independent jobs in parallel

61

```
D1 = load 'data1' ...
```

```
D2 = load 'data2' ...
```

```
D3 = load 'data3' ...
```

```
C1 = join D1 by a, D2 by b
```

```
C2 = join D1 by c, D3 by d
```

C1 and C2 are two independent jobs that can run in parallel

# More Links



- Getting Started:
  - <http://pig.apache.org/docs/r0.15.0/start.html>
- Basics
  - <http://pig.apache.org/docs/r0.15.0/basic.html>
- Built In Functions
  - <http://pig.apache.org/docs/r0.15.0/func.html>
- User-Defined Functions
  - <http://pig.apache.org/docs/r0.15.0/udf.html>
- Control Structures
  - <http://pig.apache.org/docs/r0.15.0/cont.html>
- Shell and Utility Commands
  - <http://pig.apache.org/docs/r0.15.0/cmds.html>